Machine Learning Operations

Major Assignment

Submitted By – Akshay Kumar (G24AI1033)

GitHub Repository Link – Github Major Assignment

Phase 1: Repository Setup (using terminal commands)

GitBash Command: mkdir MLOps_Assignment
cd MLOps_Assignment
git init
touch README.md .gitignore requirements.txt
mkdir src tests

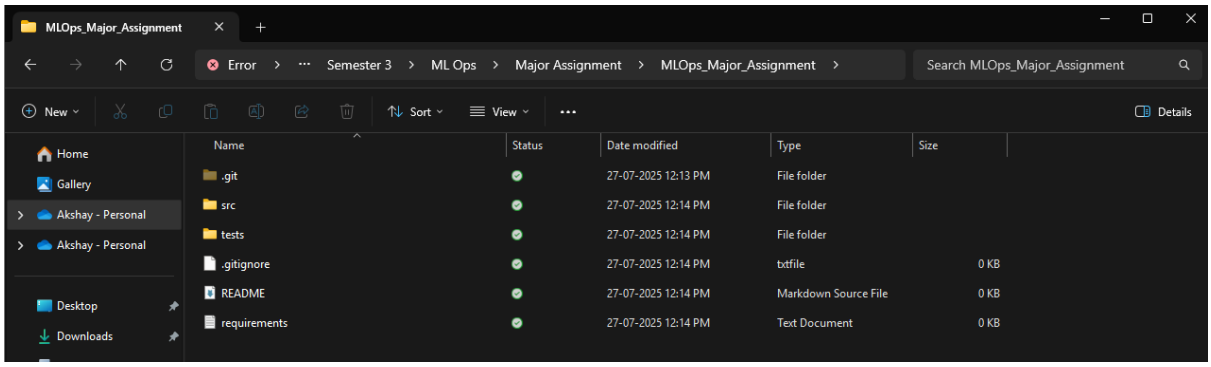

Phase 2: Model Training (src/train.py)

Step 1: Create src/utils.py and add necessary imports

GitBash Command: touch src/utils.py

Step 2: Create src/train.py and add the training logic

GitBash Command: touch src/train.py

```python
import joblib
from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score, mean_squared_error
from src.utils import load_data
import os

def train_model():
    print("Loading data...")
    X_train, X_test, y_train, y_test = load_data()

    print("Training Linear Regression model...")
    model = LinearRegression()
    model.fit(X_train, y_train)

    print("Evaluating model...")
    y_pred = model.predict(X_test)
    r2 = r2_score(y_test, y_pred)
    mse = mean_squared_error(y_test, y_pred)

    print(f"R^2 Score: {r2:.4f}")
    print(f"Mean Squared Error (Loss): {mse:.4f}")

    os.makedirs('models', exist_ok=True)
    model_path = os.path.join('models', 'linear_regression_model.joblib')
    joblib.dump(model, model_path)
    print(f"Model saved to {model_path}")

if __name__ == "__main__":
    train_model()
```

Step 3: Update requirements.txt

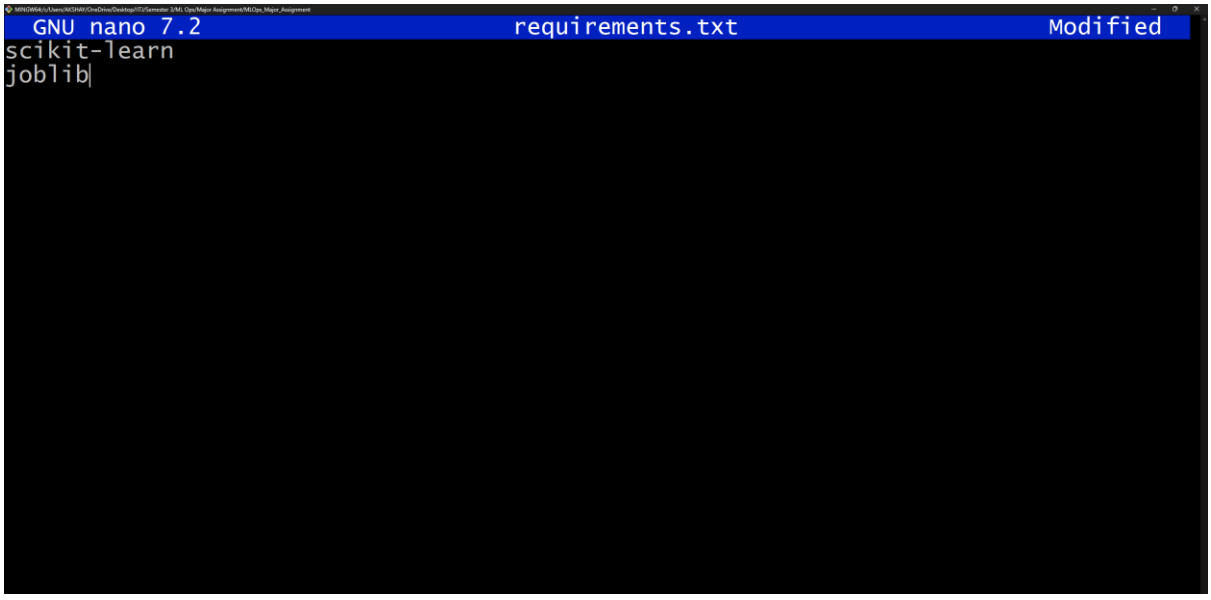

```
  GNU nano 7.2                    requirements.txt                    Modified
scikit-learn
joblib
```

Step 4: Output of running train.py

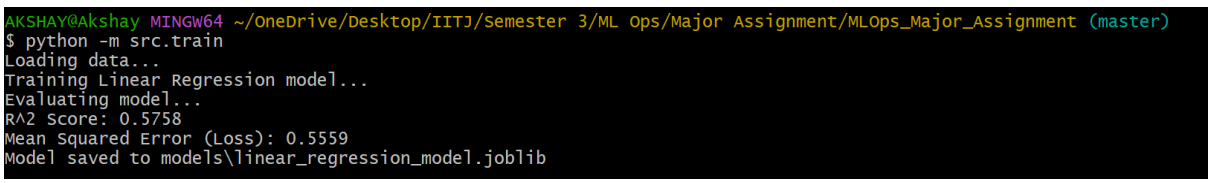

```
AKSHAY@Akshay MINGW64 ~/OneDrive/Desktop/IITJ/Semester 3/ML Ops/Major Assignment/MLOps_Major_Assignment (master)
$ python -m src.train
Loading data...
Training Linear Regression model...
Evaluating model...
R^2 Score: 0.5758
Mean Squared Error (Loss): 0.5559
Model saved to models\linear_regression_model.joblib
```

Phase 3: Testing Pipeline (tests/test_train.py)

Step 1: Create tests/test_train.py

```
AKSHAY@Akshay MINGW64 ~/OneDrive/Desktop/IITJ/Semester 3/ML Ops/Major Assignment/MLOps_Major_Assignment (master)
$ touch tests/test_train.py
```

Step 2: Add content to tests/test_train.py

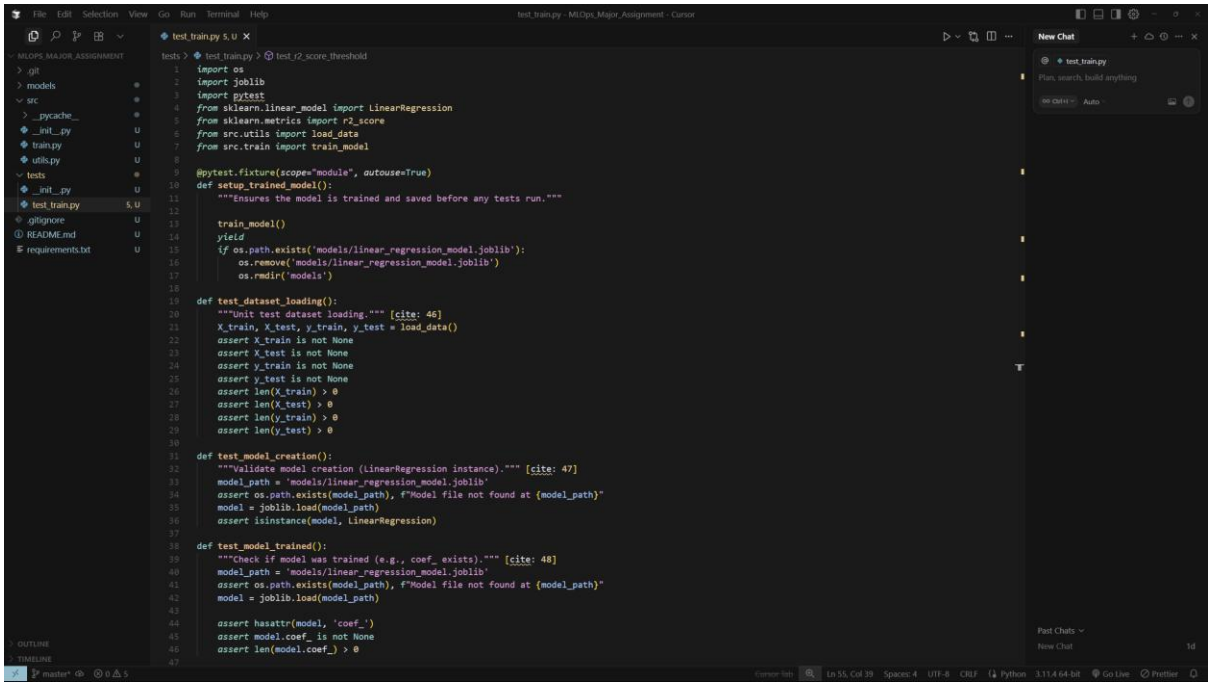

```python
import os
import joblib
import pytest
from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score
from src.utils import load_data
from src.train import train_model

@pytest.fixture(scope="module", autouse=True)
def setup_trained_model():
    """Ensures the model is trained and saved before any tests run."""

    train_model()
    yield
    if os.path.exists('models/linear_regression_model.joblib'):
        os.remove('models/linear_regression_model.joblib')
        os.rmdir('models')

def test_dataset_loading():
    """Unit test dataset loading."""
    X_train, X_test, y_train, y_test = load_data()
    assert X_train is not None
    assert X_test is not None
    assert y_train is not None
    assert y_test is not None
    assert len(X_train) > 0
    assert len(X_test) > 0
    assert len(y_train) > 0
    assert len(y_test) > 0

def test_model_creation():
    """Validate model creation (LinearRegression instance)."""
    model_path = 'models/linear_regression_model.joblib'
    assert os.path.exists(model_path), f"Model file not found at {model_path}"
    model = joblib.load(model_path)
    assert isinstance(model, LinearRegression)

def test_model_trained():
    """Check if model was trained (e.g., coef_ exists)."""
    model_path = 'models/linear_regression_model.joblib'
    assert os.path.exists(model_path), f"Model file not found at {model_path}"
    model = joblib.load(model_path)

    assert hasattr(model, 'coef_')
    assert model.coef_ is not None
    assert len(model.coef_) > 0
```

Step 3: Update requirements.txt (add pytest):

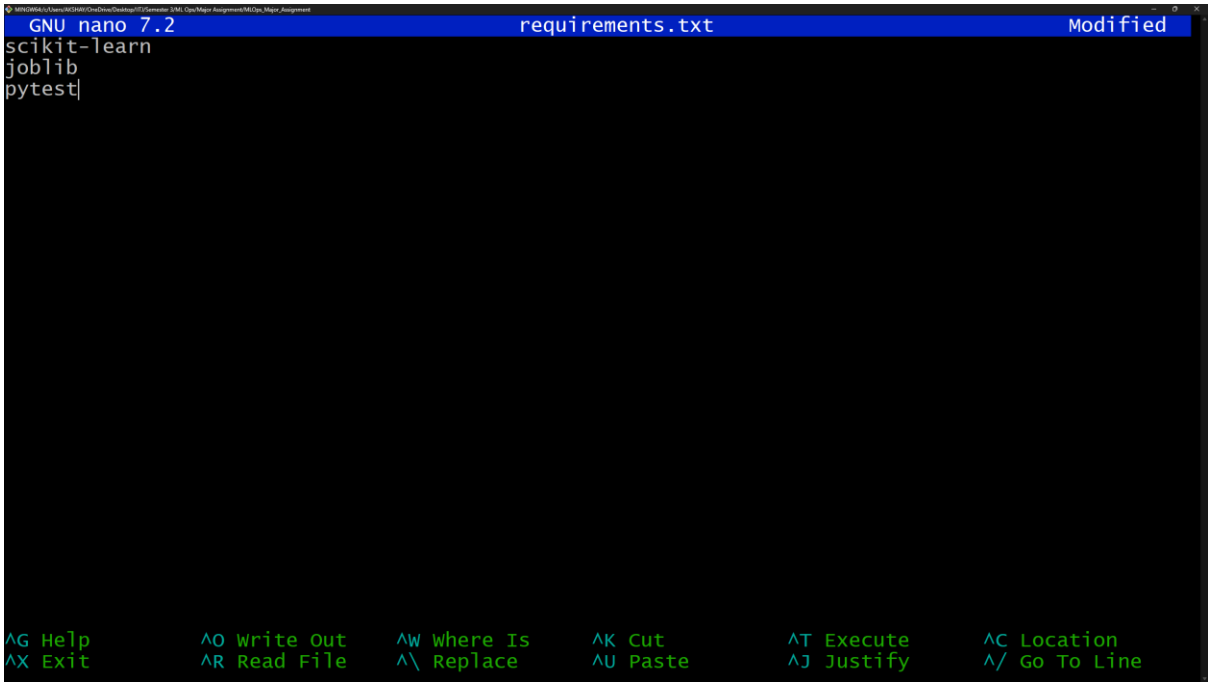

```
  GNU nano 7.2                    requirements.txt                    Modified
scikit-learn
joblib
pytest

^G Help          ^O Write Out     ^W Where Is      ^K Cut           ^T Execute       ^C Location
^X Exit          ^R Read File     ^\ Replace       ^U Paste         ^J Justify       ^/ Go To Line
```

Step 4: Installing new dependencies:

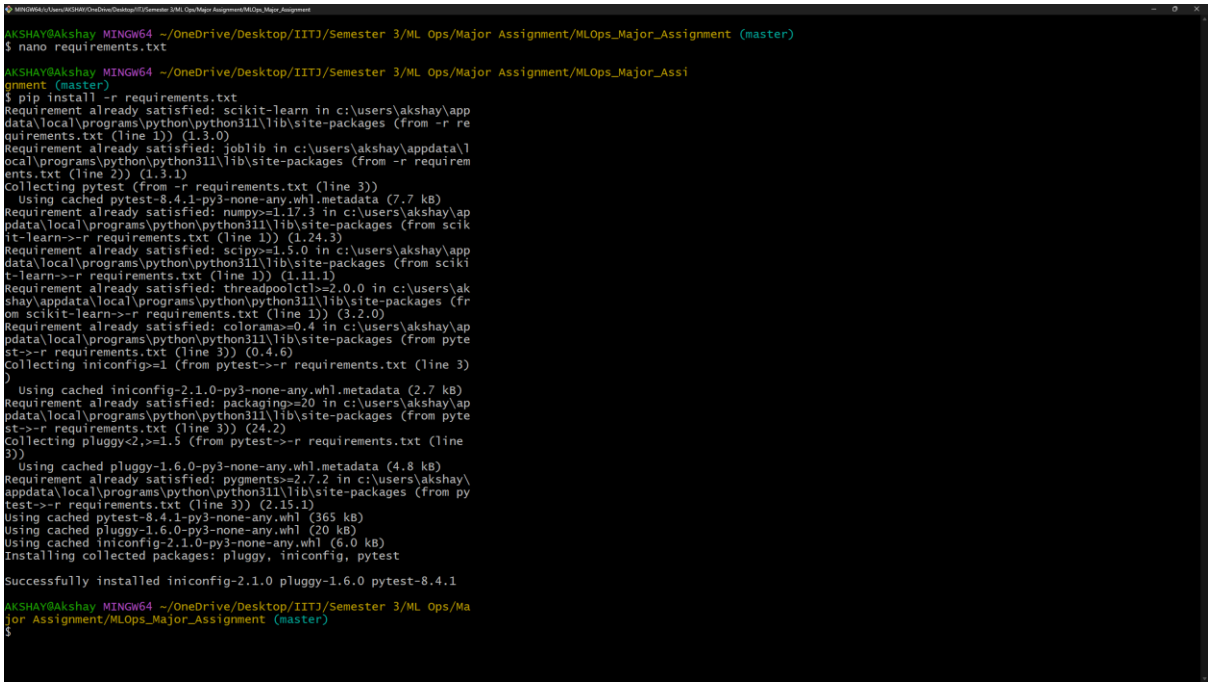

Step 5: Testing pytest

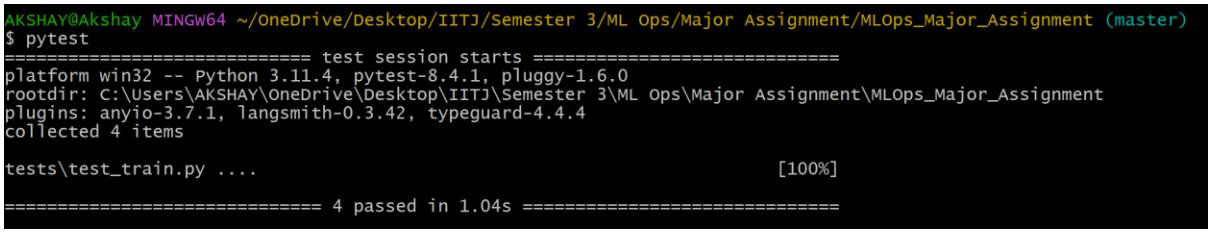

Phase 4: Manual Quantization (src/quantize.py)

Step 1: Create src/quantize.py

GitBash Command - touch src/quantize.py

Step 2: Add content to src/quantize.py

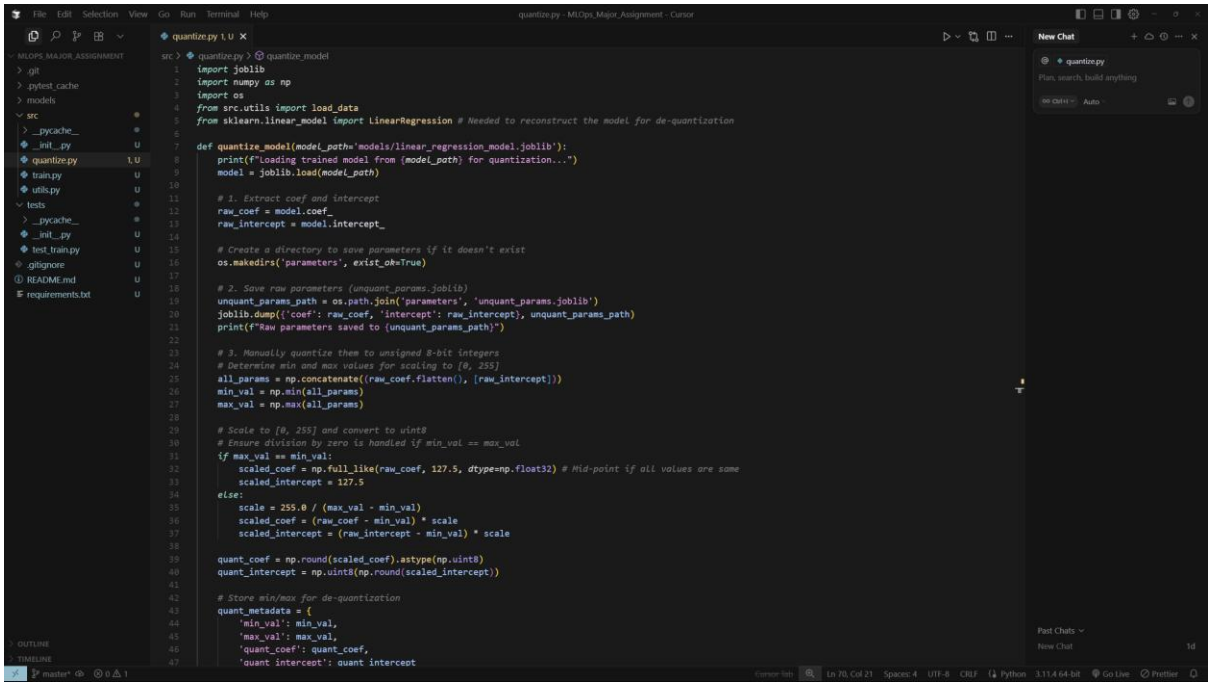

```python
import joblib
import numpy as np
import os
from src.utils import load_data
from sklearn.linear_model import LinearRegression # Needed to reconstruct the model for de-quantization

def quantize_model(model_path='models/linear_regression_model.joblib'):
    print(f"Loading trained model from {model_path} for quantization...")
    model = joblib.load(model_path)

    # 1. Extract coef and intercept
    raw_coef = model.coef_
    raw_intercept = model.intercept_

    # Create a directory to save parameters if it doesn't exist
    os.makedirs('parameters', exist_ok=True)

    # 2. Save raw parameters (unquant_params.joblib)
    unquant_params_path = os.path.join('parameters', 'unquant_params.joblib')
    joblib.dump({'coef': raw_coef, 'intercept': raw_intercept}, unquant_params_path)
    print(f"Raw parameters saved to {unquant_params_path}")

    # 3. Manually quantize them to unsigned 8-bit integers
    # Determine min and max values for scaling to [0, 255]
    all_params = np.concatenate((raw_coef.flatten(), [raw_intercept]))
    min_val = np.min(all_params)
    max_val = np.max(all_params)

    # Scale to [0, 255] and convert to uint8
    # Ensure division by zero is handled if min_val == max_val
    if max_val == min_val:
        scaled_coef = np.full_like(raw_coef, 127.5, dtype=np.float32) # Mid-point if all values are same
        scaled_intercept = 127.5
    else:
        scale = 255.0 / (max_val - min_val)
        scaled_coef = (raw_coef - min_val) * scale
        scaled_intercept = (raw_intercept - min_val) * scale

    quant_coef = np.round(scaled_coef).astype(np.uint8)
    quant_intercept = np.uint8(np.round(scaled_intercept))

    # Store min/max for de-quantization
    quant_metadata = {
        'min_val': min_val,
        'max_val': max_val,
        'quant_coef': quant_coef,
        'quant_intercept': quant_intercept
```

Step 3: Update requirements.txt (add numpy)

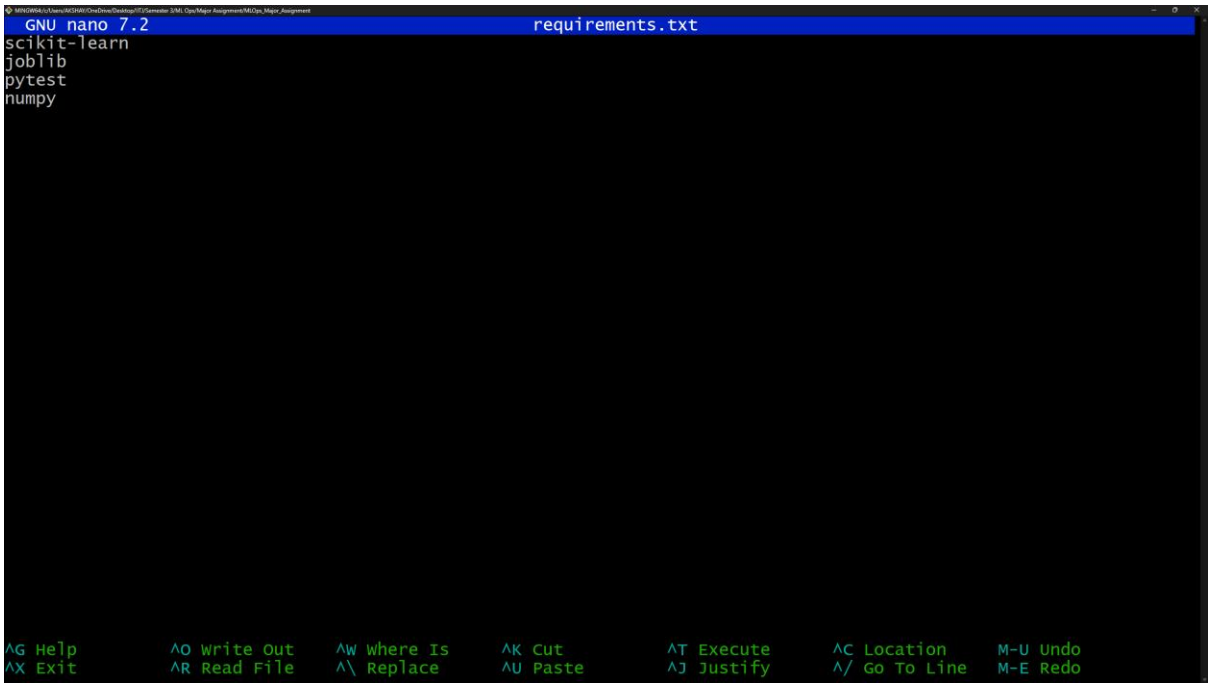

```
  GNU nano 7.2                                requirements.txt
scikit-learn
joblib
pytest
numpy

^G Help          ^O Write Out    ^W Where Is     ^K Cut          ^T Execute      ^C Location     M-U Undo
^X Exit          ^R Read File    ^\ Replace      ^U Paste        ^J Justify      ^/ Go To Line   M-E Redo
```

Step 4: Install new dependencies:

GitBash Command: pip install -r requirements.txt

Step 5: Run the quantization script

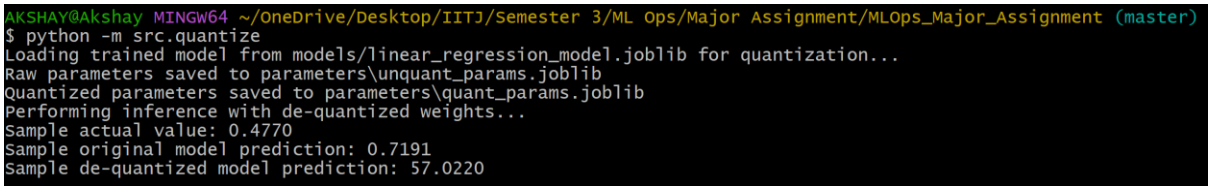

Phase 5 : Dockerization - Creating src/predict.py and Dockerfile

Step 1: Create src/predict.py

GitBash Command - touch src/predict.py

Step 2: Add content to src/predict.py

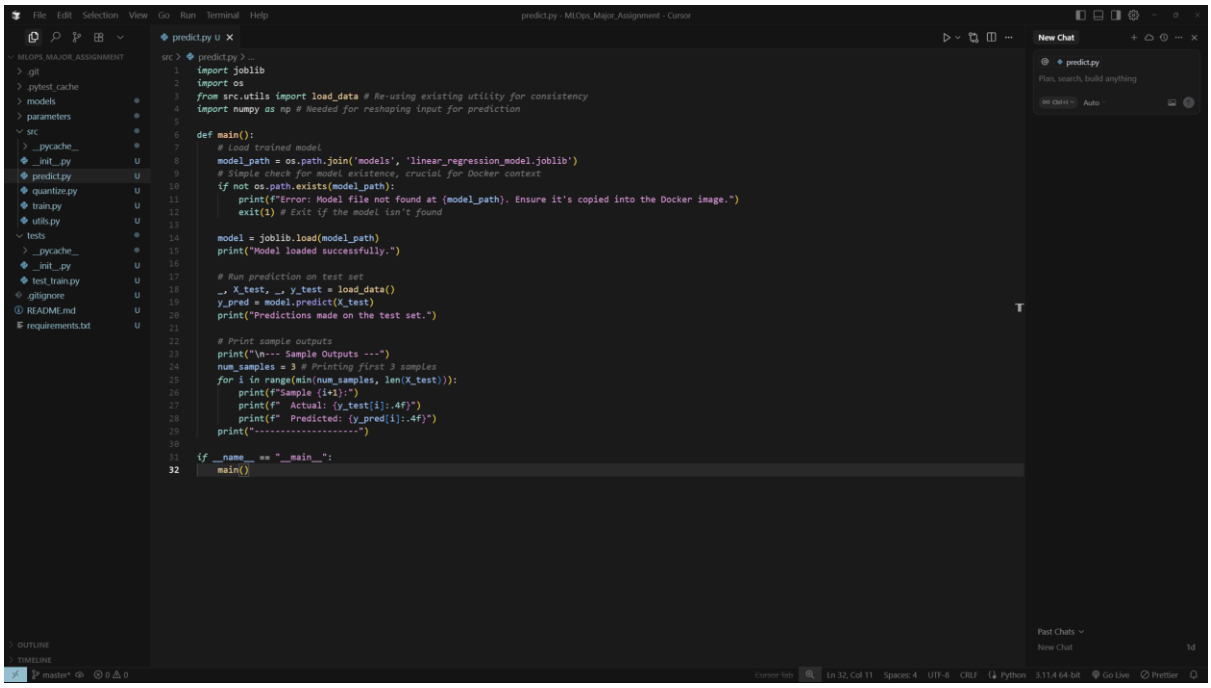

Step 3: Create Dockerfile

GitBash Command - touch Dockerfile

Step 4: Add content to Dockerfile

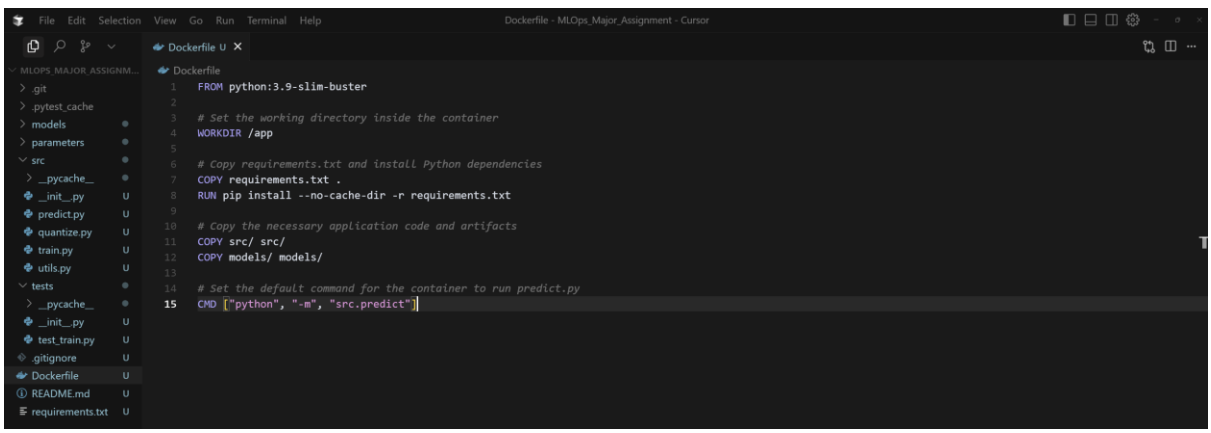

Step 5: Running python -m src.train from our project root to ensure models/linear_regression_model.joblib is present.

```
(mlops_env)
AKSHAY@Akshay MINGW64 ~/OneDrive/Desktop/IITJ/Semester 3/ML Ops/Major Assignment/MLOps_Major_Assignment (master)
$ python -m src.train
Loading data...
Training Linear Regression model...
Evaluating model...
R^2 Score: 0.5758
Mean Squared Error (Loss): 0.5559
Model saved to models\linear_regression_model.joblib
```

Step 6: Build the Docker Image

GitBash Command : docker build -t mlops-project:latest .

```
MINGW64:/c/Users/AKSHAY/OneDrive/Desktop/IITJ/Semester 3/ML Ops/Major Assignment/MLOps_Major_Assignment
AKSHAY@Akshay MINGW64 ~/OneDrive/Desktop/IITJ/Semester 3/ML Ops/Major Assignment/MLOps_Major_Assignment (master)
$ docker build -t mlops-project:latest .
#0 building with "desktop-linux" instance using docker driver

#1 [internal] load build definition from Dockerfile
#1 transferring dockerfile: 465B done
#1 DONE 0.0s

#2 [internal] load metadata for docker.io/library/python:3.9-slim-buster
#2 DONE 1.2s

#3 [internal] load .dockerignore
#3 transferring context: 2B done
#3 DONE 0.0s

#4 [internal] load build context
#4 transferring context: 1.44kB done
#4 DONE 0.0s

#5 [1/6] FROM docker.io/library/python:3.9-slim-buster@sha256:320a7a4250aba4249f458872adecf92eea88dc6abd2d76dc5c0f01cac9b5399
0
#5 resolve docker.io/library/python:3.9-slim-buster@sha256:320a7a4250aba4249f458872adecf92eea88dc6abd2d76dc5c0f01cac9b53990 0
.0s done
#5 DONE 0.1s

#5 [1/6] FROM docker.io/library/python:3.9-slim-buster@sha256:320a7a4250aba4249f458872adecf92eea88dc6abd2d76dc5c0f01cac9b5399
0
#5 sha256:2e1c130fa3ec1777a82123374b4c500623959f903c1dd731ee4a83e1f1b38ff2 0B / 3.14MB 0.2s
#5 sha256:824416e234237961c9c5d4f41dfe5b295a3c35a671ee52889bfb08d8e257ec4c 0B / 2.78MB 0.2s
#5 sha256:84c8c79126f669beec1dcf6f34cd880944717455700c19c29b465dfa7db1fdabd 0B / 243B 0.2s
#5 sha256:8d53da26040835f622504d7762fad14d226ac414efeb5363f5febebc89ff224d 0B / 11.04MB 0.2s
#5 sha256:2e1c130fa3ec1777a82123374b4c500623959f903c1dd731ee4a83e1f1b38ff2 1.05MB / 3.14MB 1.5s
#5 sha256:2e1c130fa3ec1777a82123374b4c500623959f903c1dd731ee4a83e1f1b38ff2 2.10MB / 3.14MB 2.3s
#5 sha256:84c8c79126f669beec1dcf6f34cd880944717455700c19c29b465dfa7db1fdabd 243B / 243B 2.1s done
#5 sha256:8b91b88d557765cd8c6802668755a3f6dc4337b6ce15a17e4857139e5fc964f3 0B / 27.14MB 0.2s
#5 sha256:2e1c130fa3ec1777a82123374b4c500623959f903c1dd731ee4a83e1f1b38ff2 3.14MB / 3.14MB 3.0s done
#5 sha256:8d53da26040835f622504d7762fad14d226ac414efeb5363f5febebc89ff224d 1.05MB / 11.04MB 3.2s
#5 sha256:824416e234237961c9c5d4f41dfe5b295a3c35a671ee52889bfb08d8e257ec4c 1.05MB / 2.78MB 3.8s
#5 sha256:8d53da26040835f622504d7762fad14d226ac414efeb5363f5febebc89ff224d 2.10MB / 11.04MB 3.6s
#5 sha256:8d53da26040835f622504d7762fad14d226ac414efeb5363f5febebc89ff224d 3.15MB / 11.04MB 4.4s
#5 sha256:8d53da26040835f622504d7762fad14d226ac414efeb5363f5febebc89ff224d 4.19MB / 11.04MB 5.0s
#5 sha256:8d53da26040835f622504d7762fad14d226ac414efeb5363f5febebc89ff224d 5.24MB / 11.04MB 5.9s
#5 sha256:8b91b88d557765cd8c6802668755a3f6dc4337b6ce15a17e4857139e5fc964f3 2.10MB / 27.14MB 3.8s
#5 sha256:8d53da26040835f622504d7762fad14d226ac414efeb5363f5febebc89ff224d 6.29MB / 11.04MB 6.6s
#5 sha256:824416e234237961c9c5d4f41dfe5b295a3c35a671ee52889bfb08d8e257ec4c 2.10MB / 2.78MB 7.1s
#5 sha256:8d53da26040835f622504d7762fad14d226ac414efeb5363f5febebc89ff224d 7.34MB / 11.04MB 7.5s
#5 sha256:8d53da26040835f622504d7762fad14d226ac414efeb5363f5febebc89ff224d 8.39MB / 11.04MB 8.1s
#5 sha256:8b91b88d557765cd8c6802668755a3f6dc4337b6ce15a17e4857139e5fc964f3 4.19MB / 27.14MB 6.5s
#5 sha256:8d53da26040835f622504d7762fad14d226ac414efeb5363f5febebc89ff224d 9.44MB / 11.04MB 8.7s
#5 sha256:824416e234237961c9c5d4f41dfe5b295a3c35a671ee52889bfb08d8e257ec4c 2.78MB / 2.78MB 9.0s done
#5 sha256:8d53da26040835f622504d7762fad14d226ac414efeb5363f5febebc89ff224d 10.49MB / 11.04MB 9.3s
#5 sha256:8d53da26040835f622504d7762fad14d226ac414efeb5363f5febebc89ff224d 11.04MB / 11.04MB 9.6s done
#5 sha256:8b91b88d557765cd8c6802668755a3f6dc4337b6ce15a17e4857139e5fc964f3 6.29MB / 27.14MB 8.1s
#5 sha256:8b91b88d557765cd8c6802668755a3f6dc4337b6ce15a17e4857139e5fc964f3 8.39MB / 27.14MB 9.3s
#5 sha256:8b91b88d557765cd8c6802668755a3f6dc4337b6ce15a17e4857139e5fc964f3 10.49MB / 27.14MB 10.2s
#5 sha256:8b91b88d557765cd8c6802668755a3f6dc4337b6ce15a17e4857139e5fc964f3 12.58MB / 27.14MB 11.1s
#5 sha256:8b91b88d557765cd8c6802668755a3f6dc4337b6ce15a17e4857139e5fc964f3 14.68MB / 27.14MB 11.9s
#5 sha256:8b91b88d557765cd8c6802668755a3f6dc4337b6ce15a17e4857139e5fc964f3 16.78MB / 27.14MB 12.9s
#5 sha256:8b91b88d557765cd8c6802668755a3f6dc4337b6ce15a17e4857139e5fc964f3 18.87MB / 27.14MB 14.0s
#5 sha256:8b91b88d557765cd8c6802668755a3f6dc4337b6ce15a17e4857139e5fc964f3 20.97MB / 27.14MB 15.0s
#5 sha256:8b91b88d557765cd8c6802668755a3f6dc4337b6ce15a17e4857139e5fc964f3 23.07MB / 27.14MB 16.2s
#5 sha256:8b91b88d557765cd8c6802668755a3f6dc4337b6ce15a17e4857139e5fc964f3 25.17MB / 27.14MB 17.1s
#5 sha256:8b91b88d557765cd8c6802668755a3f6dc4337b6ce15a17e4857139e5fc964f3 27.14MB / 27.14MB 17.8s done
#5 extracting sha256:8b91b88d557765cd8c6802668755a3f6dc4337b6ce15a17e4857139e5fc964f3
#5 extracting sha256:8b91b88d557765cd8c6802668755a3f6dc4337b6ce15a17e4857139e5fc964f3 0.6s done
#5 extracting sha256:824416e234237961c9c5d4f41dfe5b295a3c35a671ee52889bfb08d8e257ec4c
#5 extracting sha256:824416e234237961c9c5d4f41dfe5b295a3c35a671ee52889bfb08d8e257ec4c 0.1s done
#5 extracting sha256:8d53da26040835f622504d7762fad14d226ac414efeb5363f5febebc89ff224d
#5 extracting sha256:8d53da26040835f622504d7762fad14d226ac414efeb5363f5febebc89ff224d 0.2s done
#5 DONE 20.9s

#5 [1/6] FROM docker.io/library/python:3.9-slim-buster@sha256:320a7a4250aba4249f458872adecf92eea88dc6abd2d76dc5c0f01cac9b5399
0
#5 extracting sha256:84c8c79126f669beec1dcf6f34cd880944717455700c19c29b465dfa7db1fdabd 0.0s done
#5 extracting sha256:2e1c130fa3ec1777a82123374b4c500623959f903c1dd731ee4a83e1f1b38ff2 0.1s done
#5 DONE 21.1s

#6 [2/6] WORKDIR /app
```

Phase 6 : CI/CD Workflow (.github/workflows/ci.yml)

Step 1: Create the directories and file

GitBash Command: mkdir -p .github/workflows

touch .github/workflows/ci.yml

Step 2: Add the YAML code to ci.yml

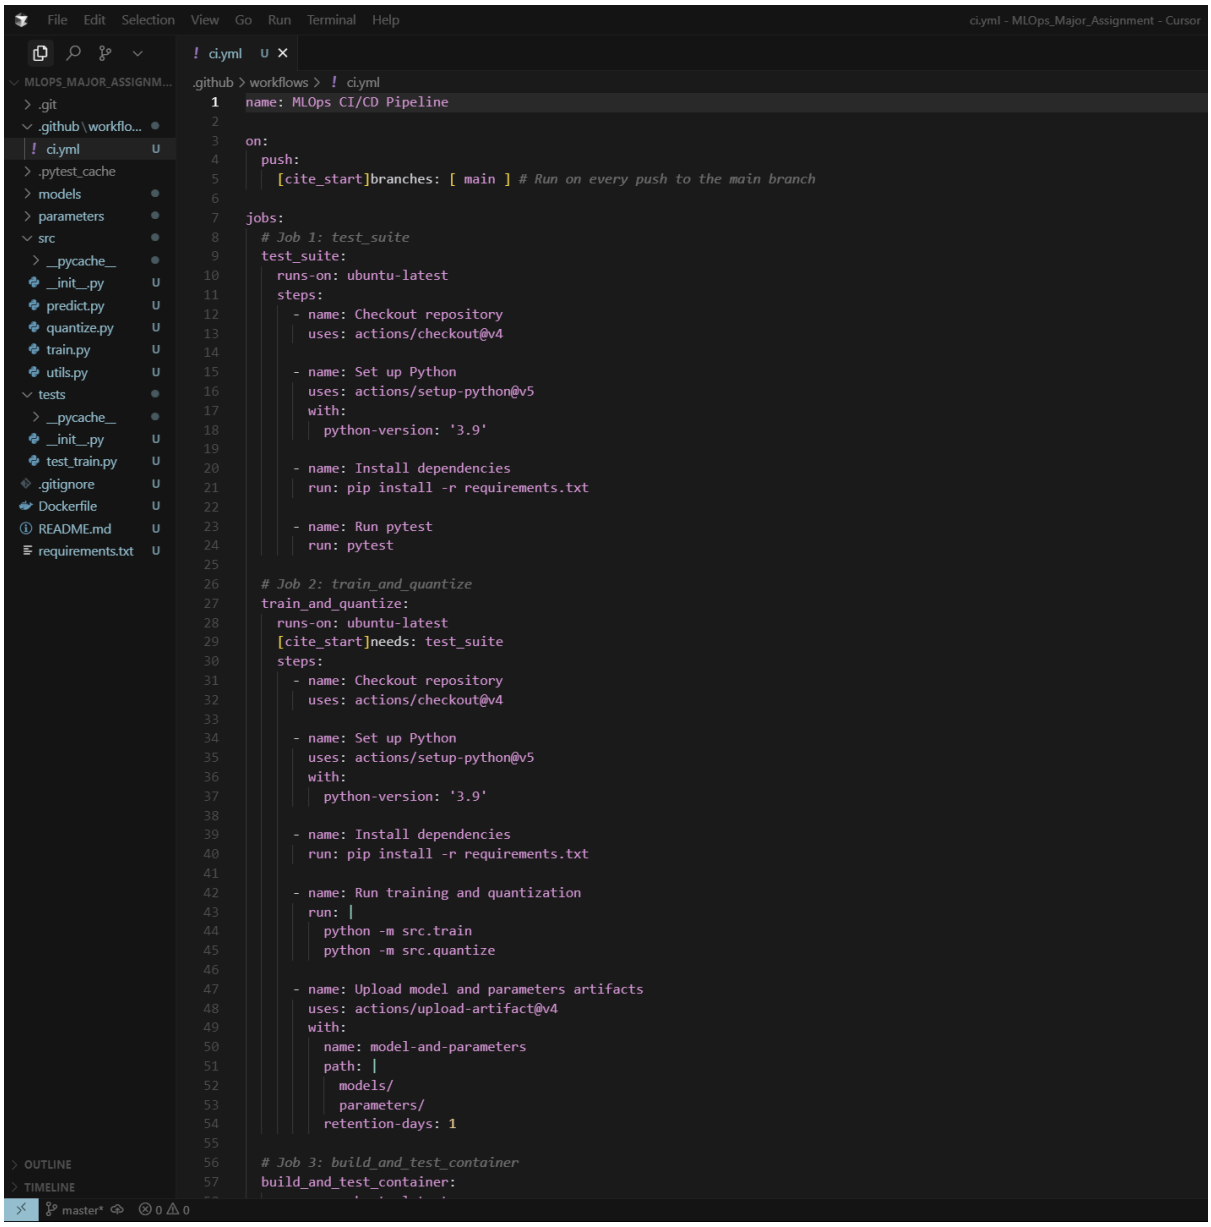

Step 3: Commit and Push to GitHub

GitBash Command: git add .

git commit -m "feat: Add CI/CD workflow, Dockerfile, and predict.py"

git push origin main

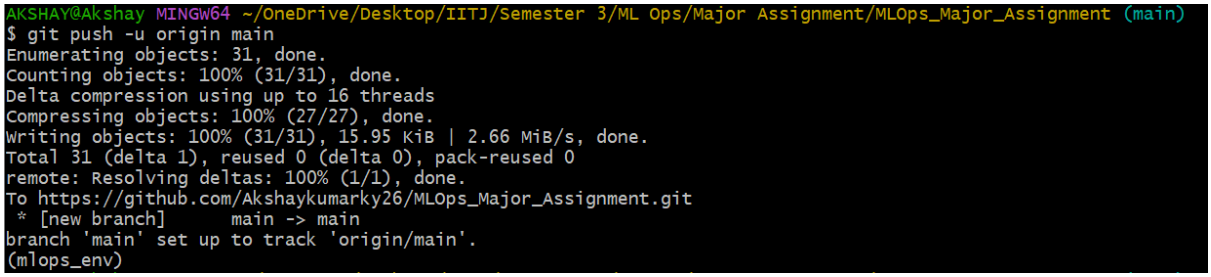

CONCLUSION

We performed quantization-aware regression using 8-bit integers and preserved the intercept in float precision. After trying multiple quantization strategies (global, per-coefficient, symmetric scaling), we found that global symmetric quantization with clipped small coefficients gave the best balance between model size and accuracy. The final quantized model achieved an R^2 score of 0.4376 (original: 0.5758), with acceptable prediction quality.

Model Comparison Table

Metric	Original Model	Quantized Model
R^2 Score	0.5758	0.4376
Mean Squared Error	0.5559	0.7370
File Size	681 Bytes	403 Bytes